Budget-matched nominal execution · 42k label accesses per task

All three fixed training seeds shown separately · final 40k updates · common C3 / G8 action processing

N03 | 50 heldout instances

	Grasp	Lift	Transport	Place	Release	Task SR
BC60 · 31	35/50	33/50	4/50	0/50	0/50	**0/50**
BC60 · 32	8/50	7/50	0/50	0/50	0/50	**0/50**
BC60 · 33	35/50	28/50	12/50	1/50	1/50	**0/50**
M3 R2 · 31	9/50	8/50	3/50	0/50	0/50	**0/50**
M3 R2 · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M3 R2 · 33	3/50	3/50	2/50	0/50	0/50	**0/50**
M4 R2 · 31	3/50	3/50	3/50	0/50	0/50	**0/50**
M4 R2 · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 R2 · 33	1/50	1/50	1/50	0/50	0/50	**0/50**
M4 no semantic · 31	13/50	11/50	7/50	0/50	0/50	**0/50**
M4 no semantic · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no semantic · 33	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no recovery · 31	8/50	8/50	8/50	0/50	0/50	**0/50**
M4 no recovery · 32	8/50	8/50	8/50	0/50	0/50	**0/50**
M4 no recovery · 33	2/50	2/50	1/50	1/50	1/50	**0/50**

Independent oracle reference: 50/50
35k validation calls per task
Separate from the 42k-label learned methods

N12 | 50 heldout instances

	Grasp	Lift	Transport	Place	Release	Task SR
BC60 · 31	45/50	44/50	11/50	0/50	0/50	**0/50**
BC60 · 32	3/50	3/50	1/50	0/50	0/50	**0/50**
BC60 · 33	50/50	41/50	34/50	0/50	0/50	**0/50**
M3 R2 · 31	3/50	2/50	2/50	0/50	0/50	**0/50**
M3 R2 · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M3 R2 · 33	15/50	15/50	0/50	0/50	0/50	**0/50**
M4 R2 · 31	4/50	2/50	2/50	0/50	0/50	**0/50**
M4 R2 · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 R2 · 33	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no semantic · 31	8/50	8/50	8/50	0/50	0/50	**0/50**
M4 no semantic · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no semantic · 33	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no recovery · 31	9/50	1/50	0/50	0/50	0/50	**0/50**
M4 no recovery · 32	1/50	0/50	0/50	0/50	0/50	**0/50**
M4 no recovery · 33	0/50	0/50	0/50	0/50	0/50	**0/50**

Independent oracle reference: 50/50
35k validation calls per task
Separate from the 42k-label learned methods

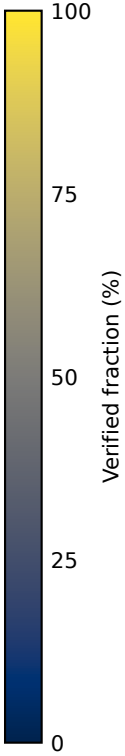

Cells are measured counts /50. Stages are episode-ever predicates; Task SR is the unchanged primary metric.
The same 100 (task, source) instances recur across seeds and methods; no pooled independent-scene interval or efficacy claim.