Budget-matched nominal execution · 28k label accesses per task

All three fixed training seeds shown separately · final 40k updates · common C3 / G8 action processing

N03 | 50 heldout instances

	Grasp	Lift	Transport	Place	Release	Task SR
BC40 · 31	27/50	25/50	1/50	0/50	0/50	**0/50**
BC40 · 32	19/50	10/50	1/50	0/50	0/50	**0/50**
BC40 · 33	35/50	30/50	18/50	0/50	0/50	**0/50**
M3 R1 · 31	10/50	8/50	8/50	0/50	0/50	**0/50**
M3 R1 · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M3 R1 · 33	2/50	2/50	0/50	0/50	0/50	**0/50**
M4 R1 · 31	6/50	6/50	6/50	1/50	1/50	**1/50**
M4 R1 · 32	5/50	5/50	5/50	2/50	1/50	**1/50**
M4 R1 · 33	4/50	4/50	2/50	0/50	0/50	**0/50**
M4 no semantic · 31	2/50	2/50	2/50	0/50	0/50	**0/50**
M4 no semantic · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no semantic · 33	2/50	2/50	0/50	0/50	0/50	**0/50**
M4 no recovery · 31	4/50	4/50	4/50	0/50	0/50	**0/50**
M4 no recovery · 32	5/50	5/50	5/50	0/50	0/50	**0/50**
M4 no recovery · 33	4/50	4/50	3/50	0/50	0/50	**0/50**

Independent oracle reference: 50/50
35k validation calls per task
Separate from the 28k-label learned methods

N12 | 50 heldout instances

	Grasp	Lift	Transport	Place	Release	Task SR
BC40 · 31	38/50	30/50	3/50	0/50	0/50	**0/50**
BC40 · 32	17/50	14/50	4/50	0/50	0/50	**0/50**
BC40 · 33	37/50	31/50	15/50	1/50	1/50	**0/50**
M3 R1 · 31	0/50	0/50	0/50	0/50	0/50	**0/50**
M3 R1 · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M3 R1 · 33	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 R1 · 31	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 R1 · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 R1 · 33	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no semantic · 31	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no semantic · 32	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no semantic · 33	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no recovery · 31	0/50	0/50	0/50	0/50	0/50	**0/50**
M4 no recovery · 32	2/50	1/50	1/50	0/50	0/50	**0/50**
M4 no recovery · 33	1/50	1/50	1/50	0/50	0/50	**0/50**

Independent oracle reference: 50/50
35k validation calls per task
Separate from the 28k-label learned methods

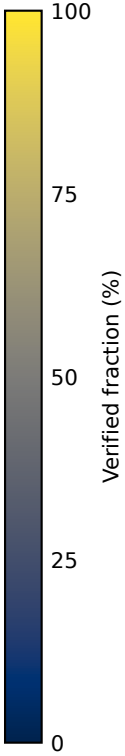

Cells are measured counts /50. Stages are episode-ever predicates; Task SR is the unchanged primary metric.
The same 100 (task, source) instances recur across seeds and methods; no pooled independent-scene interval or efficacy claim.